# Introduction

## Why Data Visualization?

Our information age more often feels like an era of information overload. Excess amounts of information are overwhelming; raw data becomes useful only when we apply methods of deriving insight from it.

Fortunately, we humans are intensely visual creatures. Few of us can detect patterns among rows of numbers, but even young children can interpret bar charts, extracting meaning from those numbers' visual representations. For that reason, data visualization is a powerful exercise. Visualizing data is the fastest way to communicate it to others.

Of course, visualizations, like words, can be used to lie, mislead, or distort the truth. But when practiced honestly and with care, the process of visualization can help us see the world in a new way, revealing unexpected patterns and trends in the otherwise hidden information around us. At its best, data visualization is expert storytelling.

More literally, visualization is a process of *mapping* information to visuals. We craft rules that interpret data and express its values as visual properties. For example, the humble bar chart in Figure 1-1 is generated from a very simple rule: larger values are *mapped* as taller bars.



*Figure 1-1. Data values mapped to visuals*

More complex visualizations are generated from datasets more complex than the sequence of numbers shown in Figure 1-1 and more complex sets of mapping rules.

## Why Write Code?

Mapping data by hand can be satisfying, yet is slow and tedious. So we usually employ the power of computation to speed things up. The increased speed enables us to work with much larger datasets of thousands or millions of values; what would have taken years of effort by hand can be mapped in a moment. Just as important, we can rapidly experiment with *alternate mappings*, tweaking our rules and seeing their output re-rendered immediately. This loop of write/render/evaluate is critical to the iterative process of refining a design.

Sets of mapping rules function as *design systems*. The human hand no longer executes the visual output; the computer does. Our human role is to conceptualize, craft, and write out the rules of the system, which is then finally executed by software.

Unfortunately, software (and computation generally) is extremely bad at understanding what, exactly, people want. (To be fair, many humans are also not good at this challenging task.) Because computers are binary systems, everything is either on or off, yes or no, this or that, there or not there. Humans are mushier, softer creatures, and the computers are not willing to meet us halfway—we must go to them. Hence the inevitable struggle of learning to write software, in which we train ourselves to communicate in the very limited and precise syntax that the computer can understand.

Yet we continue to write code because seeing our visual creations come to life is so rewarding. We practice data visualization because it is exciting to see what has never before been seen. It is like summoning a magical, visual genie out of an inscrutable data bottle.

## Why Interactive?

Static visualizations can offer only precomposed "views" of data, so multiple static views are often needed to present a variety of perspectives on the same information. The number of dimensions of data is limited, too, when all visual elements must be present on the same surface at the same time. Representing multidimensional datasets fairly in static images is notoriously difficult. A fixed image is ideal when alternate views are neither needed nor desired, and required when publishing to a static medium, such as print.

Dynamic, interactive visualizations can empower people to explore the data for themselves. The basic functions of most interactive visualization tools have changed little since 1996, when Ben Shneiderman of the University of Maryland first proposed a

"Visual Information-Seeking Mantra": *overview first, zoom and filter, then details-on-demand.*

This design pattern is found in most interactive visualizations today. The combination of functions is successful, because it makes the data accessible to different audiences, from those who are merely browsing or exploring the dataset to those who approach the visualization with a specific question in search of an answer. An interactive visualization that offers an overview of the data alongside tools for "drilling down" into the details may successfully fulfill many roles at once, addressing the different concerns of different audiences, be they those new to the subject matter or already deeply familiar with the data.

Of course, interactivity can also encourage engagement with the data in ways that static images cannot. With animated transitions and well-crafted interfaces, some visualizations can make exploring data feel more like playing a game. Interactive visualization can be a great medium for engaging an audience who might not otherwise care about the topic or data at hand.

## Why on the Web?

Visualizations aren't truly visual unless they are *seen*. Getting your work out there for others to see is critical, and publishing on the web is the quickest way to reach a global audience. Working with web-standard technologies means that your work can be seen and experienced by anyone using a recent web browser, regardless of the operating system (Windows, Mac, Linux, Android, iOS) or device type (laptop, desktop, smartphone, tablet).

Best of all, everything covered in this book can be done with freely accessible tools, so the only investment required is your time. And everything we'll talk about uses open source, web-standard technologies.

By avoiding proprietary software and plug-ins, you can ensure that your projects are accessible on the widest possible range of devices, from typical desktop computers to tablets and even phones. The more accessible your visualization, the greater your audience and your impact. In the past few years, D3's (and Mike Bostock's) insistence on supporting only web standards has helped move the web forward dramatically; very few organizations now would argue for a Flash-based chart or interactive map.

## What This Book Is

This book is a practical introduction to merging three practices—data visualization, interactive design, and web development—using D3, a powerful tool for custom, web-based visualization.

These chapters grew out of my own process of learning how to use D3. Many people, including myself, come to D3 with backgrounds in design, mapping, and data visualization, but not programming and computer science.

D3 has a bit of an unfair reputation for being hard to learn. D3 itself is not so complicated, but it operates in the domain of the web, and the web *is* complicated. Using D3 comfortably requires some prior knowledge of the web technologies with which it interacts, such as HTML, CSS, JavaScript, and SVG. Many people (myself included) are self-taught when it comes to web skills. This is great, because the barrier to entry is so low, but problematic because it means we probably didn't learn each of these technologies from the ground up—more often, we just hack something together until it seems to work, and call it a day. Yet successful use of D3 requires understanding some of these technologies in a fundamental way.

Because D3 is written in JavaScript, learning to use D3 often means learning a lot about JavaScript. For many data vis folks, D3 *is* their introduction to JavaScript (or even web development generally). It's hard enough to learn a new programming language, let alone a new tool built on that language. D3 will enable you to do great things with JavaScript that you never would have even attempted. The time you spend learning both the language and the tool will provide an incredible payoff.

My goal is to reduce that learning time, so you can start creating amazing projects sooner. We'll take a ground-up approach, starting with the fundamental concepts and gradually adding complexity. I don't intend to show you how to make specific kinds of visualizations so much as to help you understand the workings of D3 well enough to take those building blocks and generate designs of your own creation.

## Who You Are

You may be an absolute beginner, someone new to data vis, web development, or both. (Welcome!) Perhaps you are a journalist interested in new ways to communicate the data you collect during reporting. Or maybe you're a designer, comfortable drawing static infographics but ready to make the leap to interactive projects on the Web. You could be an artist, interested in generative, data-based art. Or a programmer, already familiar with JavaScript and the web, but excited to learn a new tool and pick up some visual design experience along the way.

Whoever you are, I hope that you:

- Have heard of this new thing called the "World Wide Web"
- Are a bit familiar with HTML, the DOM, and CSS
- Might even have a little programming experience already
- Aren't scared by unknown initialisms like CSV, SVG, or JSON

- Want to make useful, interactive visualizations

If any of those things are unknown or unclear, don't fear. You might just want to spend more time with Chapter 3, which covers what you really need to know before diving into D3.

# What This Book Is Not

That said, this is not intended as a traditional computer science textbook (although since the initial release of the first edition, it has been adopted by many computer science courses), and it is not intended to teach the intricacies of any one web technology (HTML, CSS, JavaScript, SVG) in depth.

In that spirit, I might gloss over some technical points, grossly oversimplifying important concepts fundamental to computer science in ways that will make true software engineers recoil. That's fine, because I'm writing for designers here. (Note to engineers: You are still welcome here!) We'll cover the basics, and then you can dive into the more complex pieces once you're comfortable.

I will deliberately *not* address every possible approach to a given problem, but will typically present what I feel is the simplest solution, or, if not the simplest, then the most understandable.

My goal is to teach you the fundamental concepts and methods of D3. Thus, this book is decidedly *not* organized around specific example projects (except for Chapter 16). Everyone's data and design needs will be different. It's up to you to integrate these concepts in the way best suited to your particular project.

# Using Sample Code

If you are a mad genius, then you can probably learn to use D3 without ever looking at any sample code files, in which case you can skip the rest of this section.

If you're still with me, you are probably still very bright but not mad, in which case you should undertake this book with the full set of accompanying code samples in hand. Before you go any further, please download the sample files from the GitHub releases page.

> This is the second edition of this book, and I've added and updated many of the code samples. On the releases page, download **version 2.0**. Different versions of the book will correspond with different releases of the code samples; I have worked very hard to ensure all the versions sync up.

Normal people will want to click the ZIP link to download a compressed ZIP archive with all the files. Hardcore geeksters will want to clone the repository using Git. If that last sentence sounds like total gibberish, please use the first option.

Within the download, you'll notice there is a folder for each chapter that has code to go with it:

```
chapter_04
chapter_05
chapter_06
chapter_07
chapter_08
…
```

Files are organized by chapter, so in Chapter 9 when I reference *01_bar_chart.html*, know that you can find that file in the corresponding location: *d3-book/chapter_9/01_bar_chart.html*.

You are welcome to copy, adapt, modify, and reuse the example code in these tutorials for any noncommercial purpose.

If you notice any errors in the examples, please report them on the GitHub issues page, and I may be able to correct them for future editions of the text.

# Thank You

Finally, this book has been handcrafted, carefully written, and pedagogically fine-tuned for maximum effect. Thank you for reading it. I hope you learn a great deal, and even have some fun along the way.